

Dict BALD: Bayesian Active Learning by Disagreement; BLR: Bayesian Linear Reg.; BN: Bayesian Network; BNN: Bayesian NN; BO: Bayesian Opt.; BP: Belief Propagation; CPD: Cond Prob Dist; DAG: Directed Acyclic Graph; DBE: Detailed Balance Eq.; DDIM: Denoising Diffusion Implicit Models; DDPG: Deep Deterministic PG; DDPM: Denoising Diffusion Prob. Models; DQN: Deep Q-Nets; ECE: Expected Calibration Error; EI: Expected Improvement; ELBO: Evidence Lower Bound; GP: Gaussian Process; HMM: Hidden Markov Model; KF: Kalman Filter; KL: Kullback-Leibler; LDM: Latent Diffusion; LOTV: Law of Total Var.; MALA: Metropolis-Adjusted Langevin; MAP: Max A Posteriori; MCMC: Markov Chain MC; MDP: Markov Decision Process; MH: Metropolis-Hastings; MI: Mutual Info; MLE: Max Likelihood Est.; MPE: Most Probable Explanation; PF: Particle Filter; PI: Prob of Improvement; POMDP: Partially Observable MDP; RBF: Radial Basis Fnc; RFF: Random Fourier Features; SGLD: Stoch Grad Langevin Dyn.; SWAG: Stoch Weight Avg Gaussian; TD: Temporal Diff.; UCB: Upper Confidence Bound; VE: Var Elimination; VI: Variational Inference; $k_{XX'} := k(X, X')$; $K_y := K_{XX} + \sigma_n^2 I$

Probability Fundamentals

Axioms: $\mathbb{P}(\Omega) = 1$; $\mathbb{P}(A) \geq 0$; Disjoint: $\mathbb{P}(\bigcup_i A_i) = \sum_i \mathbb{P}(A_i)$ **Product:** $\mathbb{P}(X_{1:n}) = \mathbb{P}(X_1) \prod_{i=2}^n \mathbb{P}(X_i | X_{1:i-1})$ **Sum:** $\mathbb{P}(X) = \sum_y \mathbb{P}(X, y)$

Bayes: $\mathbb{P}(X|Y) = \frac{\mathbb{P}(Y|X)\mathbb{P}(X)}{\mathbb{P}(Y)}$ **Cond Indep:** $X \perp Y | Z \leftrightarrow \mathbb{P}(X, Y | Z) = \mathbb{P}(X|Z)\mathbb{P}(Y|Z)$

Gaussian $\mathcal{N}(\mu, \Sigma)$: $\mathcal{N}(x; \mu, \Sigma) = \frac{1}{\sqrt{(2\pi)^d \det(\Sigma)}} \exp(-\frac{1}{2}(x - \mu)^\top \Sigma^{-1}(x - \mu))$ **Marginal:** $X_A \sim \mathcal{N}(\mu_A, \Sigma_{AA})$ **Conditional:** $X_A | X_B \sim \mathcal{N}(\mu_{A|B}, \Sigma_{A|B})$, $\mu_{A|B} = \mu_A + \Sigma_{AB}\Sigma_{BB}^{-1}(x_B - \mu_B)$, $\Sigma_{A|B} = \Sigma_{AA} - \Sigma_{AB}\Sigma_{BB}^{-1}\Sigma_{BA}$ **Linear:** $Y = MX \sim \mathcal{N}(M\mu, M\Sigma M^\top)$ **Sum:** indep $X + X' \sim \mathcal{N}(\mu + \mu', \Sigma + \Sigma')$

E., Var, Cov: $\mathbb{E}[AX + b] = A\mathbb{E}[X] + b$; **Tower:** $\mathbb{E}_Y[\mathbb{E}_X[X|Y]] = \mathbb{E}[X]$, $\text{Var}[X] = \mathbb{E}[(X - \mathbb{E}[X])^2]$; $\text{Cov}[X, Y] = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y]$, $\text{Var}[X + Y] = \text{Var}[X] + \text{Var}[Y] + 2\text{Cov}[X, Y]$ **LOTV:** $\text{Var}[X] = \mathbb{E}[\text{Var}[X|Y]] + \text{Var}[\mathbb{E}[X|Y]]$

Info Theory: **Entropy:** $H[p] = -\mathbb{E}_p[\log p(x)]$; **Gauss** $H = \frac{1}{2} \log((2\pi e)^d \det \Sigma)$ **KL:** $\text{KL}(p||q) = \mathbb{E}_p[\log \frac{p}{q}] \geq 0$; need $\text{supp}(q) \subseteq \text{supp}(p)$ **Forward KL**($p||q$): mean-seeking 覆盖; **Reverse KL**($q||p$): mode-seeking 过 confident **MI**: $I(X; Y) = H[X] - H[X|Y] = H[Y] - H[Y|X] \geq 0$, symmetric **Cond MI**: $I(X; Y|Z) = H[X|Z] - H[X|Y, Z]$ **Gauss MI**: $I[X; Y] = \frac{1}{2} \log \det(I + \sigma_n^{-2}\Sigma)$ for $Y = X + \varepsilon$ **MLE**: $\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} \sum_i \log p(y_i | x_i, \theta)$ **MAP**: $\hat{\theta}_{\text{MAP}} = \arg \min_{\theta} -\log p(\theta) + \ell_{\text{null}}$ Gaussian prior \rightarrow L2; Laplace prior \rightarrow L1

Bayesian Linear Regression

Model: $y = w^\top x + \varepsilon$, $\varepsilon \sim \mathcal{N}(0, \sigma_n^2)$; **Prior:** $w \sim \mathcal{N}(0, \sigma_p^2 I)$, L_2 正则 / weight decay; **Posterior:** $w|X, y \sim \mathcal{N}(\mu, \Sigma)$ where $\Sigma^{-1} = \sigma_n^{-2} X^\top X + \sigma_p^{-2} I$; $\mu = \sigma_n^{-2} \Sigma X^\top y$, Σ 只依赖 X , 不依赖 y **Prediction:** $y^*|x^*, X, y \sim \mathcal{N}(x^* \mu, x^* \Sigma x^* + \sigma_n^2)$ $\mu \Leftrightarrow \text{RidgeReg}$ 解(=MAP解), Σ 则 对应 其 Hessian 的逆. MAP=Ridge with $\lambda = \frac{\sigma_n^2}{\sigma_p^2}$; Online update: $O(nd^2)$

Gaussian Processes

Def: $f \sim \mathcal{GP}(\mu, k)$: any finite subset jointly Gaussian. $f_A \sim \mathcal{N}(\mu_A, K_{AA})$, $[K_{AA}]_{ij} = k(x_i, x_j)$

GP Regression: $y \sim \mathcal{N}(0, K_{XX} + \sigma_n^2 I) = \mathcal{N}(0, K_y)$

Mean: $\mu^*(x) = k(x, X)K_y^{-1}y$ **Cov:** $k^*(x, x') =$

$k(x, x') - k(x, X)K_y^{-1}k(X, x')$ **Predictive:** $y^* \sim \mathcal{N}(\mu^*(x), \sigma_n^2)$

Kernels: **Linear:** $k(x, x') = x^\top x' + \sigma_0^2$ **RBF:** $k = \exp\left(-\frac{\|x - x'\|^2}{2\ell^2}\right)$ smooth 无限 可微 **Exponential:** $k = \exp(-\|x - x'\|/\ell)$ rough **Matérn:** $\nu = 0.5 \rightarrow \text{Exp}$, $\nu \rightarrow \infty \rightarrow \text{RBF}$, $\nu \frac{1}{2}$ 制 smoothness **Periodic:** $k = \sigma^2 \exp\left(-\frac{2}{\ell^2} \sin^2\left(\frac{\pi \|x - x'\|}{\ell}\right)\right)$ **Closure:** $k_1 + k_2, k_1 \cdot k_2, c \cdot k$, $\exp(k)$ 仍 valid kernel **Stationary:** $k(x, x') = k(x - x')$; **Isotropic:** $k = k(\|x - x'\|)$

Marginal Lik: $\log p(y|X) = -\frac{1}{2}y^\top K_y^{-1}y - \frac{1}{2} \log \det(K_y) + C$ **Balance:** Data fit(前) vs Complexity(后)

Approx $O(n^3) \rightarrow \text{lower}$: **RFF:** $k(x - x') \approx \varphi(x)^\top \varphi(x')$, $O(nm^2 + m^3)$ Bochner: stationary kernel \leftrightarrow Fourier of non-neg measure **Inducing Pts:** subset $m \ll n$ points for approx

Variational Inference

Goal: Approx $p(\theta|D)$ with $q(\theta|\lambda)$ by min $\text{KL}(q||p)$

ELBO: $\mathcal{L} = \mathbb{E}_q[\log p(y|\theta)] - \text{KL}(q(\theta)||p(\theta))$ 且 $\log p(y) = \mathcal{L} + \text{KL}(q||p(\cdot|D)) \geq \mathcal{L}$ Max **ELBO** \leftrightarrow Min **KL** to posterior **Derivation:** Jensen: $\log \mathbb{E}_q[\frac{p}{q}] \geq \mathbb{E}_q[\log \frac{p}{q}]$

KL of Gaussian $\text{KL}(\mathcal{N}_p || \mathcal{N}_q) = \frac{1}{2} [\text{tr}(\Sigma_q^{-1} \Sigma_p) + (\mu_p - \mu_q)^\top \Sigma_q^{-1} (\mu_p - \mu_q) - d + \log \det(\Sigma_q^{-1} - \Sigma_p)]$

$\text{KL}(\mathcal{N}_p || \mathcal{N}_q) = \frac{1}{2} \left[\frac{(\mu_p - \mu_q)^2}{\sigma_q^2} + \frac{\sigma_p^2}{\sigma_q^2} - 1 - \log \frac{\sigma_p^2}{\sigma_q^2} \right]$ 1dim 时. **Product:** $\text{KL}(Q_X Q_Y || P_X P_Y) = \text{KL}(Q_X || P_X) + \text{KL}(Q_Y || P_Y)$

Reparam Trick: $\theta = g(\varepsilon; \lambda)$, $\varepsilon \sim \varphi$ $\mathbb{E}_{\theta \sim q}[f(\theta)] = \mathbb{E}_\varepsilon[f(g(\varepsilon; \lambda))]$ **Gaussian:** $\theta = \mu + \sigma \odot \varepsilon$, $\varepsilon \sim \mathcal{N}(0, I)$ Enable gradient: $\nabla_\lambda \mathbb{E}_q[f] = \mathbb{E}_\varphi[\nabla_\lambda f(g(\varepsilon; \lambda))]$

Laplace Approx: $q(\theta) = \mathcal{N}(\hat{\theta}, \Lambda^{-1})$ $\hat{\theta} = \text{MAP}$; $\Lambda = -\nabla^2 \log p(\hat{\theta}|D)$ (Hessian) Good at mode, overconfident elsewhere

Markov Chains & MCMC

MC basics: **Markov:** $X_{t+1} \perp X_{1:t-1} | X_t$ **Stationary** π : $\pi = \pi P$ **Irreducible:** all states reachable from any state **Aperiodic:** $\gcd\{t : P^{t(x,x)} > 0\} = 1$ **Ergodic**=Irreducible+Aperiodic: unique $\pi > 0$, $\lim_{t \rightarrow \infty} q_t = \pi$

DBE: $\pi(x)P(x'|x) = \pi(x')P(x|x')$ If satisfied \rightarrow π stationary, chain **reversible** **Proof:** sum over x' 得 $\pi(x) = \sum_{x'} \pi(x')P(x|x')$

Ergodic Thrm: $\frac{1}{n} \sum_{i=1}^n f(X_i) \xrightarrow{\text{a.s.}} \mathbb{E}_{x \sim \pi}[f(x)]$ Hoeffding: error prob decays $\exp(-n)$

MH Algorithm: Propose $x' \sim R(x'|x)$. Accept w.p.: $\alpha(x'|x) = \min\left\{1, \frac{q(x')R(x|x')}{q(x)R(x'|x)}\right\}$ Stationary: $p(x) \propto q(x)$ (unnormalized OK) Satisfies DBE \rightarrow correct stationary dist

Gibbs Sampling: Iterate: $x_i^{(t+1)} \sim p(x_i | x_{-i}^{(t)})$ Special MH with acceptance=1 **Practical:** 顺序 scan all vars, sample each from conditional

Langevin & SGLD: **Langevin:** $R(x'|x) = \mathcal{N}(x'; x - \eta \nabla f(x), 2\eta I)$ where $p \propto e^{-f}$ **MALA:** MH-corrected Langevin, poly-time for log-concave **SGLD:** $\theta_{t+1} = \theta_t + \varepsilon_t [\nabla \log p(\theta) + \nabla \log p(D|\theta)] + \sqrt{2\varepsilon_t} \xi$ Converge: $\sum_t \varepsilon_t = \infty$, $\sum_t \varepsilon_t^2 < \infty$; 常用 $\varepsilon_t \in \Theta(t^{-\frac{1}{3}})$

Gibbs Distribution: $p(x) = \frac{1}{Z} \exp(-f(x))$, f =energy function Posterior always interpretable as Gibbs

Bayesian Neural Networks

Model: Prior: $\theta \sim \mathcal{N}(0, \sigma_p^2 I)$ **Homoscedastic:** $y|x, \theta \sim \mathcal{N}(f(x; \theta), \sigma^2)$ fixed noise **Heteroscedastic:** $y \sim \mathcal{N}(f_\mu(x; \theta), \exp\{f_\sigma(x; \theta)\})$ input-dependent noise

Hetero NLL: $-\log p(y|x, \theta) = C + \frac{1}{2} [\log \sigma^2(x) + \frac{(y - \mu(x))^2}{\sigma^2(x)}]$ Model can “blame” noise but pays $\log \sigma$ penalty $\frac{1}{2}$ collapse

MAP for BNN: $\hat{\theta}_{\text{MAP}} = \arg \min \frac{1}{2\sigma_p^2} \|\theta\|^2 + \frac{1}{2\sigma_n^2} \sum_i [y_i - f(x_i; \theta)]^2$ Weight decay = Gaussian prior

Prediction: $p(y^*|x^*, D) \approx \frac{1}{m} \sum_{j=1}^m p(y^*|x^*, \theta^{(j)})$, $\theta^{(j)} \sim q$ MC approx of posterior predictive

Uncertainty Decomp: **Total Var=Aleatoric+Epistemic** Aleatoric(data noise): $\frac{1}{m} \sum_j \sigma^2(x^*, \theta^{(j)})$ Epistemic(model uncertainty): $\frac{1}{m} \sum_j [\mu(x^*, \theta^{(j)}) - \bar{\mu}]^2$ where $\bar{\mu} = \frac{1}{m} \sum_j \mu(x^*, \theta^{(j)})$

MC Dropout: $q_j(\theta_j) = p\delta_0(\theta_j) + (1-p)\delta_{\lambda_j}(\theta_j)$ Test 时 keep dropout \rightarrow 多次 forward passes \rightarrow uncertainty estimates

SWAG: Store running avg of SGD iterates: μ, Σ Space: $O(d^2)$ covariance vs $O(Td)$ all models

Calibration: Goal: Confidence \approx Accuracy **ECE**: $\sum \frac{|B_m|}{n} |\text{acc}(B_m) - \text{conf}(B_m)|$ **Temp Scaling:** $\frac{z}{T}$ on logits; $T > 1 \rightarrow$ less confident

Active Learning

Objective: $I(S) = I(f_S; y_S) = H[f_S] - H[f_S|y_S]$ NP-hard; Greedy gives $(1 - \frac{1}{e})$ -approx (submodular, monotone)

Strategies: **Uncertainty Sampling:** $x = \arg \max H[y_x|D]$ Cannot distinguish aleatoric vs epistemic **BALD:** $x = \arg \max I(\theta; y_x|D) = H[y_x|D] - \mathbb{E}_\theta[H[y_x|\theta]]$ Finds where models **disagree** about y_x **Hetero:** $x = \arg \max \{\sigma_{\text{epistemic}}^2 / \sigma_{\text{aleatoric}}^2\}$

Submodular: $F(A \cup \{x\}) - F(A) \geq F(B \cup \{x\}) - F(B)$ for $A \subseteq B$ Diminishing returns; MI is submodular

Bayesian Optimization

Regret: $R_T = \sum_{t=1}^T (f^* - f(x_t))$ Goal: sublinear $R_T/T \rightarrow 0$

Acquisition Fns: **UCB:** $x_{t+1} = \arg \max [\mu_{t(x)} + \beta_t \sigma_{t(x)}]$ $\beta_t = 0$: pure exploit; $\beta_t \rightarrow \infty$: uncertainty sampling Regret: $R_T = O(\sqrt{T\gamma_T})$ **PI:** $\text{PI}(x) = \Phi\left(\frac{\mu(x) - f(x^+)}{\sigma(x)}\right)$ **EI:** $\text{EI}(x) = (\mu - f^+) \Phi(Z) + \sigma \varphi(Z)$, $Z = \frac{\mu - f^+}{\sigma}$ **Thompson:** Sample $\tilde{f} \sim p(f|D_t)$, pick $\arg \max \tilde{f}$

Info Gain γ_T : Linear: $\gamma_T = O(d \log T)$ RBF: $\gamma_T = O((\log T)^{d+1})$ Matérn($\nu > \frac{1}{2}$): $\gamma_T = O(T^{\frac{d}{2\nu+d}} (\log T)^{\frac{2\nu}{2\nu+d}})$

MDP & RL Foundations

MDP: $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$ **Value:** $V^{\pi(s)} = \mathbb{E}\left[\sum_{t \geq 0} \gamma^t R_t | s_0 = s, \pi\right]$ **Q-fnc:** $Q^{\pi(s,a)} = R(s, a) + \gamma \sum_{s'} P(s'|s, a) V^{\pi(s')}$

Bellman Eqs: **Expectation:** $V^{\pi(s)} = R(s, \pi(s)) + \gamma \sum_{s'} P(s'|s, \pi(s))V^{\pi(s')}$ **Optimality:** $V^*(s) = \max_a [R(s, a) + \gamma \sum_{s'} P(s'|s, a)V^*(s')]$ **Q*(s, a) =** $R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^*(s', a')$ **Matrix:** $v^\pi = (I - \gamma P^\pi)^{-1} r^\pi$

Bellman's Thrm: π^* optimal iff greedy w.r.t. own V^π : $\pi^*(s) = \arg \max_a Q^*(s, a)$

PI & VI: Policy Iter: (1)Eval V^π exactly(solve LSE), (2) $\pi \rightarrow$ greedy. Fewer iters, $O(n^3)/\text{iter}$. **Value Iter:** $V \rightarrow \max_a [r + \gamma PV]$. More iters, $O(n^2m)/\text{iter}$. Both converge to optimal; VI gives ε -optimal

POMDP: Belief-state MDP: reward $\rho(b, a) = \mathbb{E}_{x \sim b}[r(x, a)]$ **Belief:** $b_t(x) = P(X_t = x | y_{1:t}, a_{1:t-1})$ **Bayes Filter:** $b_{t+1}(x) \propto o(y_{t+1}|x) \sum_{x'} P(x|x', a_t) b_t(x')$

Tabular RL

Model-based: $\hat{P}(x'|x, a) = \frac{N(x'|x, a)}{N(a|x)}$, $\hat{r}(x, a) = \text{avg rewards}$ Converges but needs many samples

Q-Learning (Off-policy): $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a') - Q(s, a))$ Uses max (ideal best a'); off-policy, model-free

SARSA (On-policy): $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$ Uses actual a' from policy; on-policy

TD Learning: $V(s) \leftarrow V(s) + \alpha(r + \gamma V(s') - V(s))$ As SGD: $\ell = \frac{1}{2}[V(s) - (r + \gamma V(s'))]^2$ Converges if Robbins-Monro: $\sum \alpha_t = \infty$, $\sum \alpha_t^2 < \infty$

Exploration: ε -greedy: prob ε random, else best **Optimistic Init:** $Q = \frac{R_{\max}}{1-\gamma}$ **Rmax:** unknown(s, a) $\rightarrow R_{\max}$, PAC guarantee

Deep RL

DQN: $\mathcal{L} = (r + \gamma \max_{a'} Q_{\theta^-}(s', a') - Q_{\theta}(s, a))^2$ Target Net θ^- : stabilize; **Experience Replay:** break correlation **Double DQN:** selection θ , eval θ^- ; reduces overestimation

Policy $\nabla_\theta J = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_t \nabla \log \pi_\theta(a_t | s_t) G_t \right]$ $\nabla \log P(\tau) = \sum_t \nabla \log \pi(a_t | s_t)$ (dynamics cancel!) **REINFORCE:** MC estimate, high variance **Baseline:** $G_t - b(s_t)$, $b = V(s)$ optimal; unbiased

Actor-Critic: **Actor:** $\pi_\theta(a|s)$; **Critic:** $V^\varphi(s)$ or $Q^\varphi(s, a)$ $\nabla J \approx \mathbb{E}[\nabla \log \pi(a|s)(Q(s, a) - V(s))]$ Critic bootstrap 减 variance 但引入 bias

Advanced: **TRPO:** $\max \mathbb{E}\left[\left(\frac{\pi_\theta}{\pi_{\text{old}}}\right) A^{\pi_{\text{old}}}\right]$ s.t. $\text{KL} \leq \delta$

DDPG: continuous actions, deterministic $\mu_\theta(s)$ **Adv Fnc:** $A^{\pi(s,a)} = Q^{\pi(s,a)} - V^{\pi(s)}$

Diffusion Models

Setup: Forward: data \rightarrow noise (fixed, no learning)
Backward: noise \rightarrow data (learned generation) Latent var model: $x_{1:T}$ are latents, x_0 is data

Forward Process: $q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1-\beta_t}x_{t-1}, \beta_t I)$ $x_t = \sqrt{1-\beta_t}x_{t-1} + \sqrt{\beta_t}\varepsilon_t$ Schedule: $\beta_t \in (0, 1)$ 单调增, $\beta_1 \approx 10^{-4}$, $\beta_T \approx 0.02$

Closed-Form Marginal \star : Define: $\alpha_t = 1 - \beta_t$, $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$ $q(x_t|x_0) = \mathcal{N}(\sqrt{\alpha_t}x_0, (1-\bar{\alpha}_t)I)$

Reparam: $x_t = \sqrt{\alpha_t}x_0 + \sqrt{1-\bar{\alpha}_t}\varepsilon$, $\varepsilon \sim \mathcal{N}(0, I)$ As $t \rightarrow T$: $\bar{\alpha}_T \rightarrow 0$, $x_T \sim \mathcal{N}(0, I)$ indep of x_0

Reverse Process: $p_{\lambda(x_{t-1}|x_t)} = \mathcal{N}(\mu_{\lambda(x_{t-1}, t)}, \sigma_t^2 I)$

Prior: $p(x_T) = \mathcal{N}(0, I)$ Generate: sample x_T , iteratively denoise to x_0

Forward Posterior: $q(x_{t-1}|x_t, x_0) = \mathcal{N}(\tilde{\mu}_t, \tilde{\beta}_t I)$
 $\tilde{\mu}_t = \frac{\sqrt{\alpha_{t-1}}\beta_t}{1-\bar{\alpha}_t}x_0 + \frac{\sqrt{\alpha_t}(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}x_t$ $\tilde{\beta}_t = \frac{(1-\bar{\alpha}_{t-1})\beta_t}{1-\bar{\alpha}_t}$ Key:
given x_0, x_t , forward posterior is Gaussian (tractable)

ELBO & Loss: $\mathcal{L} = \text{const} - \sum_{t=2}^T \underbrace{\text{KL}(q(x_{t-1}|x_t, x_0) \| p_{\lambda(x_{t-1}|x_t)})}_{L_t}$ Two Gaussians
same var: $\text{KL} \propto \|\mu_1 - \mu_2\|^2$

\star **Noise Prediction** : Predict ε instead of μ (more stable): From $x_t = \sqrt{\alpha_t}x_0 + \sqrt{1-\bar{\alpha}_t}\varepsilon$:
 $\tilde{\mu}_t = \frac{1}{\sqrt{\alpha_t}}\left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}}\varepsilon\right)$ **Simple Loss:** $L_{\text{simple}} = \mathbb{E}_{t, x_0, \varepsilon} [\|\varepsilon - \varepsilon_{\lambda(x_t, t)}\|^2]$

Training Algo: Repeat: sample $x_0 \sim p_{\text{data}}$, $t \sim \text{Unif}\{1, \dots, T\}$, $\varepsilon \sim \mathcal{N}(0, I)$ $x_t = \sqrt{\alpha_t}x_0 + \sqrt{1-\bar{\alpha}_t}\varepsilon$
 $\nabla_{\lambda} \|\varepsilon - \varepsilon_{\lambda(x_t, t)}\|^2$

Sampling Algo: $x_T \sim \mathcal{N}(0, I)$ For $t = T, \dots, 1$:
 $z \sim \mathcal{N}(0, I)$ if $t > 1$ else $z = 0$ $x_{t-1} = \frac{1}{\sqrt{\alpha_t}}\left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}}\varepsilon_{\lambda(x_t, t)}\right) + \sigma_t z$

Connection: $\varepsilon_{\lambda(x_t, t)} \approx -\sqrt{1-\bar{\alpha}_t} \nabla_{x_t} \log q(x_t)$ **De-noising = Score matching**

Variants: LDM: diffusion in VAE latent space, more efficient

DDIM: deterministic sampling, fewer steps

Cond Gen: $\varepsilon_{\lambda(x_t, t, c)}$, Classifier-Free Guidance: $\tilde{\varepsilon} = (1+w)\varepsilon_{\lambda(x_t, t, c)} - w\varepsilon_{\lambda(x_t, t)}$

QuickCheck:

- **VI:** Approx posterior via ELBO. Laplace MAP, Reparam for grad.
- **MCMC:** Sample posterior. MH accept/reject, Gibbs coordinate, Langevin uses ∇ .
- **GP:** Prior over fncs, closed-form posterior. RBF smooth, Matérn tunable.
- **BNN:** Prior on weights, MC predictive. Aleatoric=data noise, Epistemic=model.
- **Active:** Max MI, BALD for disagreement, submodular \rightarrow greedy($1 - \frac{1}{e}$).
- **BO:** UCB balance explore/exploit, EI expected gain, Thompson sample.
- **BN:** DAG factorization, d-sep for indep, BP exact on trees.
- **KF:** Linear Gaussian, Kalman gain trades predict vs observe.

• **Diffusion:** Forward=noise, Backward=denoise, train predict ε .

• **On/Off:** On=SARSA,REINFORCE,PPO; Off=Q-learn,DQN,SAC

• **Bellman:** $V = R + \gamma PV$;